


Why selective data journaling?

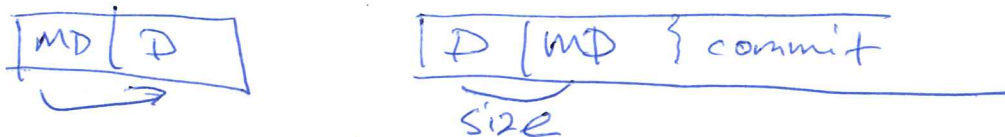
Ordered j. - semantics

Block overwrite? NO Atomicity



- data block first
- crash before journal
- mismatch btwn meta data + data


data journaling




- ① crash before commit
- ② crash after commit - new data see overwrite w/ new size

selective d.j. - perf. usually of ordered
- data journaling if detect overwrite

optimistic w/o care:




matches?



data journaling semantics?

1) reuse after notification



2) selective dij.

- overwrite block?

