

System Technique: Backpointers

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be **added** in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new **challenges** does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

LFS Bradley, Akila, Sudarshan

+

Uses backpointers to keep track of live blocks throughout the file system. The imap stores the location of the data block.

Wisckey Sudarsan, Yash, Shunmiao

Uses backpointers from value log to key for garbage collection

AFS Server has pointers to traverse directory structure.
Local system stores, ~~callbacks~~ ~~node values~~ pointers locally to be able to quickly traverse the tree and improve performance by reducing server load.

(kawshik-C, Peter)

System Technique: Background Work

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe **how** your chosen system uses the technique.
 - Explain **why** the technique is **useful** in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be **added** in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new **challenges** does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

+ For wisecry log, two pointers head and tail are maintained. wisecry validates which of the key-value pair is valid by referring the key tree and as part of garbage collection moves it to head of the log. This process is performed in background.

Advantages: ① Does not interfere with foreground.

② Can be fine tuned with respect to

'+' ↪ foreground, so that it foreground requires more resources, background

→ use of version num in inode (lfs) can be limited to limit garbage collection overhead.

VENKATESH + ++ checkpointing in journaling fs .

+ ++ ADN in OptFS .

Akila,
Sudarshan,
Bradley.

SSD Shumiao, Sudarshan, Yash

Garbage collection of SSD aims to

1) pre-erase the blocks that are dangling to improve

the write performance
2) compaction of data blocks in the background to reduce write amplification

3) wear-leveling

LFS Garbage Collection

System Technique: Layering

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one system that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

Coda's Venus presents the same interface when disconnected from the file servers. This simplifies app development, because developers don't need to worry about which mode Coda is currently in. Similarly, users of the file system can treat it as if it were a local FS.

Akhil +
 Flashgraph → Flashgraph uses SAFS to cache the data and to overlap I/O with computation. SAFS buffers edge lists from SSDs, so that Flashgraph can adapt to different cache hit rates.

Sudarshan, Akilar, Bradley. { RAID → Depending on the type of RAID used, the lowest layer can have 'n' disks. But the top layer (File system) finds it as one logical disk. So there is clear separation of layers, where FS is built on top of RAID.

System Technique: Indirection

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one system that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

System: LFS (Peter & Kausik. C) ++

- Fenced checkpoint region contains i-map pointers
i-map ~~points~~ points to inodes.
- Allowed sequential writes to i-map instead of in-place updates which would have resulted in random writes and reduced performance.

System: Wickety Tim B. Kausik. Kausik

- Wickety separately stores key in LSM-tree while value in logs. However to ease the Garbage Collection, it also stores the corresponding key along with the value in the log.
- Thus, it avoids scanning of the whole LSM-tree.
- By separating keys & values it reduces the read & write amplification,

APG
fid

SSDs - FTLs - indirection

FFS - inodes - file names

System Technique: Checksums/Parity

optf

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

FS Review
group G

X Optimistic crash calculate checksum using both metadata and data.

S By computing a checksum over the entire transaction and placing its value in J_c , the writes to J_m and J_c can be issued together. Improving performance. (crash recovery can avoid replay of improperly committed transactions.) \rightarrow slightly unclear

✓ Fs - Review 8 Minh, Nikhita, Bidgit RAID 4 - single parity disk

RAID 5 - Rotating parity

+ RAID 4 uses parity to provide Reliability, good sequential Performance + RAID 5 uses parity and parallelism to give better Seq and Random throughput

+ RAID 5 uses parity and parallelism to give better Seq and Random throughput

* Performance in terms of CPU usage goes down. As extra compute is required to compute checksums.

WiskKey \rightarrow Can benefit from addition of checksums. Possible because key and values are at different places. checksum can bring them to a consistent state

Disk: Checksums

key, ptr, checksum
over
data

System Technique: Optimism

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe **how** your chosen system uses the technique.
 - Explain **why** the technique is **useful** in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be **added** in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new **challenges** does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

Yash
Sudarshan
Shunmiao

✓ ① Optimistic journaling
Opt FS → separate ordering & durability ; sacrifice durability for performance . Benefit : higher performance .

Coda
+

2) assume that files won't be overwritten by multiple users ; write conflicts are unlikely.
This allows coda to implement reliability for disconnected systems.
(- Peter E. Kaushik - C)

However - when conflicts do happen , you need handle that

System Technique: Sorting

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one system that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

(Aarati Kakaraparthy) Graph Chi: We have a list of ordered vertices, which we partition.

Zhenwen Song Kevin Houch In each shard, we sort the edges by source. The benefit of doing this is, when out-edges corresponding to a shard are accessed in other shards, they are sequential. + +

~~Log Structured File System~~ LFS

+ (If there are workloads which write random block and then read sequential reads then Log structured file system doesn't give performance.) (For these kind of workloads we can sort the blocks before writing.) → LFS does not do this, not sure how this would work

↑
Sudeep
Song
Leon

Fs Review 8
Wisc Key

→ uses sorting of keys in the LSM tree

+ The technique of sorting keys is essential for fast searching and by separating keys and values and storing keys in memory it is optimized for SSDs

System Technique: More work done more efficiently

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one system that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

1.

In FlashGraph, the authors were able to do more work by focusing on vertex paths than edges. *X load balancing*
They used a thread based model; partitioned across vertices.
As soon as a thread was done, it started stole others' work.
The key here is to load balance in an efficient way, by being flexible to re-assign partitions to different workers. [Surya, Karan, Brandon]

→ isn't clear.
→ (they don't do "more" work)

2. Whiskey: Uses prefetching for range queries.

(Enclita, Pukkit) It leverages parallel I/O characteristics of SSDs to prefetch values from the vlog concurrently based on the access pattern of the range query. (∴ values are not sorted) '+'

3. X-Stream: It's edge-centric. It scans the edge list for multiple times

(Zhenwen Song) sequentially. It reads more than required, but this work is sequential, which is more efficient than the random access in SSDs or HDDs.
way

System Technique: Overlap/Parallelism

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe how your chosen system uses the technique.
 - Explain why the technique is useful in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be added in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new challenges does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

+ Flash Graph: Group 8: Minh; Nikhita; Bidyut
+ Overlapping: FG hides latency by overlapping computation with I/O

(X) + Parallelism: FG initiates many parallel I/Os and process data when it's ready

+ Why it's useful: + get maximum performance out of SSD.
→ Out-of-memory graph computation faster on SSD

RAID : Group 9 [Brandon, Surya, Kanan]

(+) Parallelism RAID allows us to read multiple blocks of data at the same time across multiple disks. It takes advantage of having multiple disks. It speeds up read rates and improves crash consistency.

SSD - Unwritten Contract:

Use of NCQ depth and channels to serve requests in parallel

01

WiscKey

System Technique: Caching

1. Read the notes written by the previous groups for this technique. Please rate their response with a "+", a checkmark, or a "-"; if you see anything you disagree with or think is wrong or unclear, please politely correct.
2. For the System Technique listed above, think about each of the papers/systems you have read in CS 736. Pick one **system** that uses that technique in an interesting or important way (that has NOT been already described on this page).
 - In a few sentences, describe **how** your chosen system uses the technique.
 - Explain **why** the technique is **useful** in that system (i.e., the benefits of the technique).

If you are unable to think of a different system that uses this technique, you can instead (in order of preference):

- Describe how that technique could be **added** in a useful way to one of the systems
- Elaborate on the answer given previously by a different group; for example, what new **challenges** does this technique introduce that must now be solved?
- Pick a different technique you haven't considered yet (indirection, layering, checksums, caching, overlap/parallelism, background work, more work done more efficiently, optimism, or sorting), denote the technique, and answer this question for that technique.

Please alternate who in your group takes notes and write your first names by your notes.

1. AFS uses whole file caching on local disks of client workstations. (+)
Advantages: (i) Reduces network traffic and server load (ii) Client only contacts server on open & close (iii) allows use of bulk data transfer protocols (iv) reduces refetch of data after reboots (v) simplifies cache management
{ Suchitav, Pulkit }

2. X-stream: Use caching of vertices in CPU cache and edges in RAM, or vertices in RAM depending on graph size. This improves processing time by speeding up random accesses by using the faster memory tier. +

3. WiscKey: The original LSM has significant read amplification, which make RAM as a cache less useful. WiscKey only has ~~xx~~ keys in the tree, which is small enough to store in RAM. With this caching, random access to keys is much faster.
{ Leon, Song, Pradeep }

- Coda
- SSDs (Overwritten) - FTL
- LFS - imap