

What are the metrics or considerations one should take into account when building a CPU scheduler?

Response time - interactive jobs

Fairness - same amount of CPU \leftarrow user process threads - org

Priority

Low overheads \leftarrow low # migrations
context switches

Utilization high -

Simplicity - easy to modify, understand, prove properties?
- easy to debug reproduce

No starvation

Avoid contention \leftarrow load balance across cores
global planning

Non-critical, bg - low priority \rightarrow deadlines - realtime

Power consumption

cache consideration - cache affinity

hw aware - heterogeneity

scalable \leftarrow multiple cores 64? - load balancing
processes - single queue

Simplicity, Manageability, Easy to test/understand
Hard to game. prove prop.
High throughput for cpu-bound jobs
Fast turnaround/response time for interactive jobs
High utilization (\rightarrow work conserving)
Few context switches
Meet deadlines - realtime
Proportional share - across users, groups, processes
No priority inversion -
Energy conservation
Cache re-use/sharing
Heterogeneity
Locks - no spinwaiting on same core.
Run @ same time as cooperating threads. $\xrightarrow{\text{signal}} \text{wait}$
Low lock contention for shared queues. $\sim \# \text{cores}$
Scalable w/ $\#$ threads.
Balanced load across cores
~~Work~~ Compose w/ other resources $\left\{ \begin{array}{l} \text{memory} \\ \text{disk} \end{array} \right.$

Utilization - Linux ~~to~~ Work Conserving
no idle resources, if waiting job

Good to be Non Work conserving?
resources idle?

- 1) power consumption - patch work
- 2) time critical deadline - no preemption
idle B - more important!
preemption costly
~~Am~~ - parallel systems

3) cache affinity - state

A \leftarrow cache-B \rightarrow B
A \leftarrow 11111111

Disks! anticipatory scheduling

A1 \leftarrow A2 \longleftrightarrow B1 \rightarrow B2

A1 i/o
A2 comp
A1 A2
A1 A2

B1
B2
B1 ... B2